

# Product Requirements Document (PRD)

## PromiseKeeper — The B2B Obligation Ledger

**Version:** 2.0 (Market-Verified Edition) **Date:** July 2026 **Status:** For Review — competitive landscape verified via live market research (see §11); final AppExchange listing-by-listing sweep still required before build kickoff **Platform:** Salesforce (100% native) + Agentforce **Distribution:** AppExchange / AgentExchange

---

### 1. Executive Summary

PromiseKeeper is a Salesforce-native, Agentforce-powered application that automatically extracts, tracks, and drives fulfillment of *mutual commitments* ("promises") made between sellers and buyers across emails, call transcripts, and meeting notes. It creates a bilateral obligation ledger on the Opportunity, Account, or Case — "you owe them 2 things; they owe you 1" — then agentically follows through: nudging internal owners, drafting external reminders for rep approval, detecting fulfillment signals, and escalating breaches before they erode trust or stall the deal.

**One-line pitch:** An agent that catches every "I'll send that by Friday" — from both sides — and makes sure it actually happens.

**Core insight:** Deals and renewals rarely die from one big failure; they die from dozens of dropped micro-commitments. No CRM today has a native construct for a *mutual obligation* extracted automatically from unstructured communication, tracked against fulfillment, and escalated on breach. PromiseKeeper makes the Promise a first-class CRM object.

**Positioning thesis (post market verification):** AI *extraction* of action items is now table stakes — Gong, Einstein Conversation Insights, AI notetakers, and even Slackbot all detect action items. What no one has built is what happens *after* extraction: a bilateral obligation object with a full lifecycle (confirmed → fulfilled/broken/renegotiated), **evidence-verified closure**, buyer-side accountability, and reliability analytics computed from kept-vs-broken history. PromiseKeeper's moat is not "we detect commitments" — it is "**we verify they were kept, for both sides, and turn the record into trust data.**" Extraction is the on-ramp; the *verified ledger* is the product.

---

### 2. Problem Statement

#### 2.1 The problem

- B2B deals, renewals, and delivery engagements generate a constant stream of micro-commitments scattered across email threads, calls, and meetings: "I'll send the security questionnaire by Friday," "We'll loop in legal next week," "I'll get you the pricing exception by EOD Thursday."

- These commitments are **bilateral** — the buyer owes things too — but no system tracks the two-sided ledger of who owes what, to whom, by when.
- Salesforce Tasks are manually created, one-sided (internal only), and rarely reflect what was actually promised in conversation.
- Conversation intelligence tools (Gong, Chorus, etc.) surface seller action items in their own UI, but do not track buyer obligations, do not live natively on the CRM record, and do not agentically drive fulfillment.
- Mutual action plan tools (Recapped, Accord) require manual plan construction and are mostly non-native.

## 2.2 The cost of the problem

- **Trust erosion:** Broken small promises compound; buyers quietly downgrade their perception of vendor reliability.
- **Cycle-time drag:** Deals stall waiting on items nobody remembered to send or chase.
- **Win-rate impact:** Sales methodology research consistently ties responsiveness and follow-through to win rates.
- **Renewal risk:** Post-sale, unfulfilled commitments (SLAs, promised features, onboarding items) directly drive churn.
- **Manager blindness:** Sales managers coach on pipeline stages, not on commitment hygiene, because the data doesn't exist.

## 2.3 Who feels it

| Persona                        | Pain today                                                                                                                 |
|--------------------------------|----------------------------------------------------------------------------------------------------------------------------|
| Account Executive (AE)         | Forgets promises made across 30+ active threads; embarrassed on calls when a buyer asks "did you ever send that?"          |
| Customer Success Manager (CSM) | Renewal conversations derailed by unfulfilled commitments from the sales cycle or delivery                                 |
| Sales Manager / RevOps         | No visibility into which deals are stalled on outstanding obligations, theirs or the buyer's                               |
| Delivery / PMO                 | Contractual and project commitments live in email, not in a system of record                                               |
| Buyer (external)               | Loses trust when vendors drop items; also forgets their own commitments (loop in legal, share data, book the exec meeting) |

## 3. Product Vision & Goals

### 3.1 Vision

Every commitment made in the course of a B2B relationship — by either side — is captured, visible, tracked, and fulfilled. The Promise becomes a standard unit of CRM data, and "promise health" becomes as fundamental a deal signal as stage or amount.

### 3.2 Goals

1. **Capture:** Automatically extract  $\geq 90\%$  of explicit commitments from connected communication channels with correct owner, beneficiary, and due date.
2. **Fulfillment:** Increase on-time fulfillment of tracked commitments by a measurable margin (target: +30% vs. baseline in pilot orgs).
3. **Outcomes:** Demonstrate correlation between promise-ledger health and win rate / cycle time / renewal rate within 2 quarters of deployment.
4. **Adoption:**  $\leq 10$  minutes to first value ("aha" moment: seeing the auto-built ledger on an existing Opportunity); daily active use by reps via pre-meeting briefs.
5. **Trust & safety:** Zero unauthorized external communication — every outbound message is human-approved by default.

### 3.3 Non-goals (v1)

- Not a full mutual action plan / deal room product (no shared external workspace in v1).
  - Not a conversation intelligence suite (no talk-track analytics, sentiment scoring of calls, coaching scorecards).
  - Not a contract lifecycle management (CLM) tool — contractual obligation tracking is a future expansion, not v1.
  - Not a task manager replacement — Promises sync to Tasks; we don't rebuild task UX.
- 

## 4. Personas & Jobs-to-be-Done

### 4.1 Primary personas

#### P1 — Account Executive "Ava"

- JTBD: "Before every meeting, tell me what I owe them and what they owe me, and never let me drop a promise."
- Success: Walks into calls prepared; closes loops without mental overhead; uses buyer obligations as legitimate re-engagement reasons.

#### P2 — Customer Success Manager "Chris"

- JTBD: "Track every commitment made during onboarding and the relationship so renewals aren't ambushed by broken promises."

- **Success:** Renewal QBRs show a clean fulfillment record; at-risk accounts flagged early by promise breaches.

### P3 — Sales Manager / RevOps "Maya"

- **JTBD:** "Show me which deals are stalled on obligations and who habitually drops commitments so I can coach and forecast better."
- **Success:** Pipeline reviews include promise-health; forecast calls flag deals blocked on buyer-side obligations.

## 4.2 Secondary personas

- **Delivery/PMO lead:** tracks post-sale delivery commitments on Cases/Projects.
  - **Sales Engineer:** owns technical follow-ups (security docs, architecture diagrams) assigned by the agent.
  - **Admin:** installs, configures channels, permissions, and agent guardrails.
  - **Executive sponsor (buyer of the app):** wants win-rate/cycle-time reporting to justify spend.
- 

## 5. Core Concepts & Data Model

### 5.1 The Promise object (custom object: `Promise__c`)

The heart of the product. A Promise is a discrete, trackable commitment extracted from communication or entered manually.

**Key fields:**

| Field                  | Type                 | Notes                                                                                          |
|------------------------|----------------------|------------------------------------------------------------------------------------------------|
| Promise Name / Summary | Text                 | Agent-generated, e.g., "Send SOC 2 report"                                                     |
| Description            | Long text            | Full context, verbatim source quote reference                                                  |
| Direction              | Picklist             | <div>We Owe Them</div> / <div>They Owe Us</div> / <div>Internal</div> / <div>Third Party</div> |
| Owner (Internal)       | Lookup(User)         | The internal person responsible (if We Owe Them / Internal)                                    |
| External Owner         | Lookup(Contact)      | The buyer-side person who made the commitment (if They Owe Us)                                 |
| Beneficiary            | Lookup(Contact/User) | Who the promise was made to                                                                    |
| Due Date               | Date/DateTime        | Extracted ("by Friday" → resolved to actual date) or inferred                                  |
| Due Date Confidence    | Picklist             | Explicit / Inferred / Unknown                                                                  |
| Status                 | Picklist             | Proposed → Confirmed → In Progress → Fulfilled / Broken / Waived / Renegotiated / Disputed     |
| Fulfillment Evidence   | Long text / links    | Link to email, file, event that satisfied it                                                   |
| Source Type            | Picklist             | Email / Call Transcript / Meeting Notes / Slack / Manual                                       |
| Source Record          | Polymorphic lookup   | EmailMessage, VoiceCall, Event, ContentDocument, etc.                                          |
| Related To             | Lookup               | Opportunity / Account / Case / Contract / custom object                                        |
| Extraction Confidence  | Percent              | Model confidence; drives review queue routing                                                  |
| Priority / Deal Impact | Picklist             | Agent-assessed: Blocking / High / Normal / Low                                                 |
| Renegotiation History  | Related list         | Prior due dates and reasons (see 5.3)                                                          |

| Field          | Type     | Notes                                                |
|----------------|----------|------------------------------------------------------|
| Verbatim Quote | Text     | The exact sentence(s) the promise was extracted from |
| Visibility     | Picklist | Internal Only / Shareable Externally                 |

### 5.2 Supporting objects

- **Promise Event** (`PromiseEvent__c`): Immutable audit log of every state change, nudge sent, snooze, renegotiation — powers analytics and dispute resolution.
- **Nudge** (`Nudge__c`): A reminder instance (internal notification or external draft) with channel, approval status, and outcome.
- **Promise Settings (custom metadata)**: Org-level and team-level configuration (channels, thresholds, quiet hours, approval policies).
- **Ledger Snapshot** (`LedgerSnapshot__c`): Periodic per-record snapshot for trend reporting (promise debt over time).

### 5.3 Promise lifecycle

```

Detected (agent) → Proposed (in review queue if low confidence)
  → Confirmed (auto if high confidence, or human-approved)
    → In Progress
      → Fulfilled (evidence attached, auto or manual)
      → Broken (due date passed + grace period, no evidence)
      → Renegotiated (new due date agreed; history preserved)
      → Waived (no longer needed; reason captured)
      → Disputed (human flags extraction as wrong → feeds model feedback
loop)

```

#### Design principles:

- Renegotiation is a first-class flow, not an edit — real-world promises get re-dated constantly, and the *history* of renegotiations is itself a health signal.
- "Broken" is never announced externally by the agent; it is an internal analytic state.
- Every Promise retains a verbatim quote + deep link to its source for instant human verification.

## 6. Functional Requirements

### 6.1 Commitment Extraction Engine (F-EXT)

#### F-EXT-1 — Multi-source ingestion

- v1 (MVP): Email (Einstein Activity Capture / Email-to-Salesforce / Gmail & Outlook integrations).
- v1.5+: Call transcripts (Salesforce Voice, Gong/Chorus/Fireflies via API), meeting notes (uploaded or synced), calendar event descriptions.
- v2: Slack channels/DMs (with explicit consent scoping), chat (LiveChat/Messaging sessions on Cases).

**F-EXT-2 — Extraction specification.** For each detected commitment, the agent must extract:

- The commitment action (normalized summary + verbatim quote)
- Who made it (internal user vs. external contact — resolved against CRM records)
- Who it benefits
- Due date, resolving relative expressions ("by Friday," "before our next call," "end of month") against the message timestamp and, where relevant, the meeting calendar
- Conditionality ("if you sign by June, we'll include onboarding free") → flagged as Conditional Promise with trigger condition
- Related record (Opportunity/Case/Account) via thread matching, contact roles, and existing activity relations

### **F-EXT-3 — Confidence-based routing**

- High confidence (configurable threshold, default  $\geq 85\%$ ): auto-create as Confirmed; rep notified.
- Medium: created as Proposed; appears in rep's daily review queue ("Did you promise this? ✓ / ✗ / Edit").
- Low: discarded but logged for tuning (admin-visible sample review).

### **F-EXT-4 — Deduplication & threading**

- The same promise repeated across a thread ("just following up on that SOC 2 report") must update the existing Promise, not create duplicates.
- Reply-detection: a message like "here it is, attached" in the same thread is a fulfillment signal, not a new promise.

### **F-EXT-5 — Negative/retraction detection**

- Detect cancellations ("actually, we no longer need the security review") → propose Waived.
- Detect renegotiations ("can we push that to next Wednesday?") → propose Renegotiated with new date.

### **F-EXT-6 — Language & tone handling**

- v1: English. v2: multilingual extraction (prioritize by market demand — DE, FR, ES, PT, JA).
- Handle hedged commitments ("I'll try to get that over this week") — captured with a Soft Commitment flag; configurable whether soft commitments are tracked.

## 6.2 Fulfillment Detection (F-FUL)

**F-FUL-1 — Signal monitoring.** The agent monitors for evidence a promise was kept:

- Attachment sent in the thread matching the promised artifact (filename/content-type/semantic match)
- Meeting booked matching a promised meeting ("I'll set up time with our architect")
- Reply from the obligated party stating completion ("just sent it over" / "legal is now looped in — adding Priya")
- CRM signals: file uploaded to the record, quote sent, contract stage change, Case status change
- New contact added matching a "loop in X" promise

**F-FUL-2 — Evidence linking.** On detection, attach evidence link(s) and set Status = Fulfilled (auto if high confidence; otherwise one-click confirm for the owner).

**F-FUL-3 — Manual fulfillment.** One-click "Mark fulfilled" with optional evidence, from record page, brief, or notification.

**F-FUL-4 — Breach detection.** Due date + configurable grace period elapsed without evidence → Status = Broken (internal) + escalation policy triggers.

## 6.3 Nudges & Follow-Through (F-NDG)

**F-NDG-1 — Internal nudges**

- Configurable cadence (e.g., T-2 days, due date, overdue daily) via Salesforce notifications, email digest, and (premium) Slack DM.
- Smart batching: one digest per user per day, not notification spam. Urgent/blocking promises can break through.
- Snooze / renegotiate / waive actions inline in the nudge.

**F-NDG-2 — External reminder drafts (premium)**

- For "They Owe Us" promises approaching or past due, the agent drafts a polite, context-aware reminder email in the rep's voice, referencing the original commitment naturally (never robotic, never shaming).
- **Human-in-the-loop by default:** drafts require rep approval before send. Org admins may enable auto-send for specific low-stakes categories only, with explicit opt-in and audit logging.



- Tone controls: friendly / neutral / firm; org-level style guidance; per-relationship overrides.
- Frequency guardrails: max external nudges per contact per week (default 1); blackout for contacts marked sensitive; respects opt-outs.

### **F-NDG-3 — Escalation policies (premium)**

- Rule builder: e.g., "If a Blocking promise we owe is >3 days overdue → notify manager"; "If buyer promise blocking stage progression is >7 days overdue → suggest mutual-plan reset talk track to rep."
- Escalations never go external without human action.

**F-NDG-4 — Suggested actions.** Where the agent can help fulfill the promise itself, it offers: draft the promised email, generate the promised summary doc, book the promised meeting (with Scheduler), or open the relevant record.

## **6.4 Pre-Meeting Briefs (F-BRF)**

**F-BRF-1** — Before any calendar event with external attendees tied to a tracked record, deliver a brief (notification + record page + optional Slack):

- "You owe them: [2 items, statuses, one-click actions]. They owe you: [1 item + suggested gentle raise]. Recently fulfilled: [items worth acknowledging — 'thanks for sending the data']."
- Include renegotiation candidates ("the API docs promise is due tomorrow; propose new date on the call?").

**F-BRF-2** — Post-meeting: within minutes of transcript/notes availability, present "New promises detected in this meeting" for one-tap confirmation.

**F-BRF-3** — Daily digest option: morning summary of all promise activity across a user's book of business.

## **6.5 Ledger Views & UX (F-UX)**

**F-UX-1 — Record-page Ledger component** (LWC on Opportunity/Account/Case/custom objects):

- Two-column bilateral view: "We owe them" / "They owe us," sorted by due date, with status chips, overdue highlighting, evidence links, and inline actions (fulfill, nudge, renegotiate, waive, dispute).
- Ledger balance header: "3 open · 1 overdue (ours) · 2 open (theirs) · 14 fulfilled all-time · 93% our on-time rate."

**F-UX-2 — My Promises workspace:** cross-record personal view for reps/CSMs: everything I owe, everything owed to my accounts, review queue, and history.

**F-UX-3 — Team command center (manager):** promise debt by rep/team/deal, breach trends, deals blocked on buyer obligations, coaching views (habitual late items by category).

**F-UX-4 — Review queue:** fast keyboard-friendly triage of Proposed promises (confirm / edit / reject) with the verbatim quote and source deep link visible.

**F-UX-5 — Utility bar & global actions:** quick-add manual promise; global search across promises.

**F-UX-6 — Mobile:** briefs and one-tap actions in Salesforce mobile app.

**F-UX-7 — Shareable ledger (v2, opt-in):** export or share a filtered, external-safe view of mutual commitments (PDF/link) for QBRs and mutual action alignment — bridging toward mutual action plans without becoming a deal room.

## **6.6 Agentforce Conversational Interface (F-AGT)**

**F-AGT-1 — Natural-language interaction via Agentforce:**

- "What do I owe Acme?" / "What's blocking the Globex renewal?" / "Draft a reminder to Priya about the legal review" / "Mark the SOC 2 promise fulfilled — I sent it this morning" / "Push everything due this week on Initech to next Monday and draft a note explaining."

**F-AGT-2 — Agent actions catalog** (invocable, reusable in customer's own agents/flows): ExtractPromises, GetLedger, CreatePromise, UpdatePromiseStatus, DraftExternalNudge, GenerateBrief, RenegotiatePromise.

**F-AGT-3 — Explainability:** every agent-created or agent-changed record answers "why?" with source quote, signal, and confidence.

## **6.7 Analytics & Reporting (F-ANL) (premium)**

- **Promise Health Score** per record: composite of open debt, overdue ratio, renegotiation churn, bilateral balance, and velocity of fulfillment.
- **Deal-health integration:** surface Promise Health as a field usable in reports, forecasts, list views, and Einstein/revenue-intelligence models; flag "stalled on buyer obligation" deals.
- **Trust analytics:** our on-time rate vs. theirs, by account, segment, rep, team, promise category (docs, meetings, pricing, legal, technical).
- **Outcome correlation dashboards:** win rate / cycle time / renewal rate vs. promise-health cohorts.
- **Org benchmarks:** anonymized category benchmarks (v2, opt-in).
- Standard report types + prebuilt CRM Analytics dashboards; all data reportable natively.

## **6.8 The Reliability Graph (F-REL) (premium — v2.0 flagship differentiator)**

The kept/broken/reneegotiated record compounds into a novel analytics layer no adjacent product can produce:

- **F-REL-1 — Reliability signals per entity:** on-time fulfillment rate, average slip days, renegotiation frequency, and responsiveness computed for each internal user, external contact, account, team, and promise category. Presented as directional signals (e.g., High / Moderate / Watch) with full drill-down to underlying promises — never as a public "score" that could shame individuals.
- **F-REL-2 — Deal application:** pre-meeting briefs surface counterpart patterns ("this stakeholder's commitments typically land ~4 days late — build buffer into the close plan"); stage-gate warnings when a deal's critical path runs through a low-reliability obligation.
- **F-REL-3 — Promise-backed forecasting:** open blocking obligations and counterpart reliability feed a "commitment-adjusted close date" signal alongside the rep's forecast date — grounding forecast calls in obligation reality rather than optimism.
- **F-REL-4 — Ethics guardrails:** external-contact reliability data is internal-only, aggregate-first, never included in outbound drafts, and excludable per org policy. The point is planning realism, not scoring people.

## 6.9 Commitment Escrow — Conditional Promises (F-ESC) (v2.0 differentiator)

Structured handling of "if X, then Y" commitments — the most expensive promise type and the least tracked anywhere:

- **F-ESC-1 — Conditional extraction:** "If you sign by the 30th, we'll include premium onboarding free" → Promise created in **Escrow** state with machine-readable trigger (signature by date) and payload (onboarding inclusion).
- **F-ESC-2 — Trigger monitoring:** agent watches for the trigger condition (stage change, contract signature, date passage); on trigger, the payload promise auto-activates with the correct owner — including cross-team activation (CSM inherits a promise the AE made).
- **F-ESC-3 — Handoff integrity:** at closed-won, all escrowed and open promises transfer visibly to the post-sale ledger; the #1 leak (sales concessions forgotten by delivery) is structurally closed.
- **F-ESC-4 — Concession registry:** managers get a per-deal view of what has been conditionally offered — informal discount/concession governance no CPQ or CI tool captures because it lives only in conversation.

## 6.10 Source-agnostic ingestion strategy (F-SRC) (anti-commoditization)

PromiseKeeper positions **above** extractors, not against them: native connectors ingest action-item outputs from Einstein Conversation Insights, Gong, Fireflies, and Slack meeting intelligence as candidate promises, then apply the lifecycle, verification, bilateral, and reliability layers those tools lack. Extraction competitors become distribution channels.

## 6.11 Administration & Setup (F-ADM)

- Guided setup: connect channels, choose objects to track, set thresholds, pick approval policy, select pilot team.
  - Permission sets: PromiseKeeper User / Manager / Admin; respects org sharing model (promises inherit visibility from their parent record by default).
  - Category taxonomy management (default set + custom categories).
  - Exclusion rules: domains, contacts, record types, or keywords never processed (e.g., HR, legal-privileged threads).
  - Model feedback console: disputed extractions, precision/recall sampling, threshold tuning.
  - Sandbox-friendly test mode with synthetic data pack for demos.
- 

## 7. Use Cases & Scenarios

### 7.1 New business (Opportunity)

1. **Classic drop-catch:** AE promises a security questionnaire "by Friday" on a Tuesday call. Promise auto-created, SE auto-suggested as owner (security category routing). Thursday nudge; Friday the SE replies with attachment; auto-fulfilled with evidence. Buyer never had to chase.
2. **Buyer-side accountability:** Buyer says "I'll get you the data schema next week." Following Wednesday, agent drafts: "Hope the week's going well — whenever you get a chance to share that data schema we discussed, we can have the integration assessment back to you within 48 hours." Rep approves, tweaks, sends. Deal keeps moving, and the nudge reads as service, not pestering.
3. **Stalled-deal diagnosis:** Manager filters pipeline by "blocked on their obligation >7 days" — instantly sees which stalled deals are actually waiting on the buyer, changing forecast conversations and re-engagement strategy.
4. **Pre-demo brief:** "You owe them pricing for the EU add-on (due today — draft attached). They owe you the list of pilot users. They fulfilled the security contact intro yesterday — worth a thank-you."
5. **Conditional promise:** "If you can sign by the 30th, we'll include premium onboarding." Tracked as conditional; if signature lands by the date, the onboarding promise activates on the Account with the CSM as owner — the classic sales-to-delivery handoff leak, closed.

### 7.2 Renewals & post-sale (Account/Case)

6. **Handoff integrity:** All unfulfilled sales-cycle promises transfer visibly to the CSM at closed-won. Nothing promised in the deal silently evaporates in delivery.

7. **Renewal armor:** At renewal QBR, CSM presents the fulfillment record: 41 commitments tracked, 39 fulfilled on time. Trust made legible.
8. **Churn early-warning:** Rising broken-promise count (ours) on an account feeds health scoring before NPS ever drops.
9. **Support promises:** Agent tracks "we'll ship the fix in the next release" made on Cases — the single most commonly broken support promise category — and alerts when the release ships without it.

### 7.3 Management & operations

10. **Coaching:** Maya sees a rep habitually renegotiates technical-doc promises → root cause is SE capacity, not rep behavior → staffing fix, not a performance conversation.
11. **Process mining:** Category analytics show "pricing exception" promises average 6 days late org-wide → RevOps fixes the desk process.
12. **Forecast hygiene:** Commit-stage deals with open blocking obligations get flagged in the forecast call.

### 7.4 Vertical scenarios

- **Professional services:** SOW-adjacent promises during scoping tracked into delivery; PMO ledger per project.
  - **Healthcare vendor management:** compliance-document promises (BAAs, certifications) with strict due-date tracking and audit trail.
  - **Government contracting:** formal Q&A and submission commitments with immutable audit log (Promise Events) supporting compliance posture.
  - **Channel/partner (future):** three-way ledgers across vendor-partner-customer.
- 

## 8. Agentforce Architecture

### 8.1 Agent topology

- **Extraction Agent (batch + event-driven):** triggered on new EmailMessage / transcript availability; runs extraction spec (6.1); writes Promises + Events. Uses Data Cloud-grounded prompts with record context (contact roles, open promises for dedup, calendar).
- **Fulfillment Monitor (event-driven):** subscribes to platform events (email received, file uploaded, event created, stage changed); performs evidence matching.
- **Nudge Orchestrator (scheduled):** evaluates cadence and escalation policies; generates internal notifications and external drafts; enforces frequency guardrails.
- **Brief Generator (scheduled/pre-event):** compiles briefs from ledger state + recent activity.

- **Conversational Agent (Agentforce topic + actions):** user-facing Q&A and commands (6.6).

## 8.2 Platform components

- Custom objects & fields (5.1–5.2); Platform Events for signal fan-in; Flow-invocable actions for all agent capabilities (admin-composable); LWCs for ledger, queue, briefs; Prompt templates in Prompt Builder (customer-tunable within guardrails); Data Cloud optional for cross-channel identity resolution and transcript storage at scale.

## 8.3 Design constraints

- **100% native:** no data leaves the Salesforce trust boundary for core processing (uses Einstein Trust Layer / Agentforce LLM gateway). This is a primary differentiator vs. non-native CI tools and must survive security review marketing claims.
  - All writes idempotent and replay-safe (email sync can deliver duplicates).
  - Governor-limit-conscious batch design for high-volume email orgs; per-org throttles.
- 

## 9. Trust, Safety, Privacy & Compliance

This product reads private communications and can contact customers. Trust design is a feature, not a checkbox.

- **Consent & scoping:** process only channels/mailboxes explicitly connected; per-user opt-in for mailbox processing where required; honor org exclusion rules (6.8).
- **Human-in-the-loop:** external sends require approval by default; auto-send is opt-in, category-scoped, audited, and revocable.
- **No external shaming:** the agent never tells a buyer they "broke" a promise; external language is always forward-looking and courteous. Hard-coded tone guardrails + red-team test suite for draft generation.
- **Privileged/sensitive content:** keyword and pattern-based skip lists (legal privilege, HR, health data as applicable); admin-auditable skip log.
- **Data protection:** GDPR/CCPA alignment — verbatim quotes are minimized excerpts; right-to-erasure flows cascade to Promise/Event records referencing a contact; configurable retention for source excerpts.
- **Security review:** target Salesforce AppExchange security review from day one; document threat model (prompt injection via inbound email content is the top LLM risk — inbound text is treated as untrusted data, never as instructions; extraction prompts are injection-hardened and outputs schema-validated).
- **Auditability:** immutable Promise Event log; every agent decision explainable with source + confidence.
- **Hallucination controls:** promises must be anchored to a verbatim source quote; no quote → no promise. Dispute flow feeds evaluation set; publish precision targets (see

---

## 9A. Market Potential Assessment (added in v2.0)

**Verdict: strong potential, contingent on execution of the verification engine and correct positioning.**

**Demand evidence (verified):** Every major vendor is investing in exactly the adjacent problem — Gong markets AI agents that flag missed commitments; Salesforce ships auto action items in ECI and Slackbot and pitches follow-through as a core value; the MAP/digital-sales-room category (Dock, Topo, SalesHood — the latter reporting large win-rate lifts from buyer-seller alignment features) exists purely because bilateral accountability sells. Heavy investment on all sides of the gap, with the gap itself (verified bilateral ledger) unoccupied, is the classic signature of a real category opening.

**Market sizing logic (directional):** Horizontal across all B2B Salesforce orgs running Sales/Service Cloud — hundreds of thousands of orgs, tens of millions of seats. Serviceable early market: mid-market and enterprise complex-sales teams (10–500 sellers/CSMs) — the segment that feels dropped-commitment pain acutely but finds Gong's price point (platform fee + ~\$1.4–1.6K/user/yr) heavy. A \$30–60/user/month wedge under Gong with a native trust story addresses a multi-billion-dollar spend pool (conversation intelligence + sales engagement + MAP categories combined) with a differentiated entry point.

**Structural tailwinds:** (1) Agentforce/AgentExchange is new and actively promoted — early quality listings get disproportionate visibility; (2) "revenue intelligence" messaging fatigue makes trust/accountability positioning fresh; (3) AI-agent-verified work products (evidence over claims) is the direction of the whole market; (4) post-sale obligation management is an untouched expansion space.

**Honest risks to potential:** platform absorption by Salesforce (highest), Gong down-market motion, and category-education cost — "obligation ledger" is a new mental model, mitigated by the 60-second self-building-ledger demo. Net: the idea clears the bar for build, with the \$11.5 sweep as the final gate.

---

## 10. Packaging, Pricing & Go-to-Market

### 10.1 Editions

|         | Free Demo  | Team (Premium)    | Enterprise              |
|---------|------------|-------------------|-------------------------|
| Sources | Email only | + Calls, meetings | + Slack, chat, custom   |
| Users   | 3          | Per-seat          | Per-seat + platform fee |

---

|           | Free Demo     | Team (Premium)                 | Enterprise                                               |
|-----------|---------------|--------------------------------|----------------------------------------------------------|
| Nudges    | Internal only | + External drafts, escalations | + Auto-send policies                                     |
| Briefs    | Daily digest  | Pre/post-meeting               | + Custom cadences                                        |
| Analytics | Ledger only   | Health score, dashboards       | + Benchmarks, API, Data Cloud                            |
| Admin     | Basic         | Policies, categories           | SSO-adjacent controls, sandbox seeding, priority support |

## 10.2 Why customers pay monthly

- Value accrues per-deal, continuously; tied to metrics (win rate, cycle time, NRR) teams already report on.
- The Promise ledger becomes irreplaceable proprietary CRM data — switching cost grows with history.
- Demo lands in 60 seconds: connect email → watch the ledger build itself on a live Opportunity.

## 10.3 Land-and-expand

Land with one sales team (free/pilot) → expand seats → attach CS/renewals → attach premium analytics → post-sale obligation tracking (future SKU).

## 10.4 Positioning

"Trust and accountability infrastructure" — deliberately fresh against "revenue intelligence" fatigue. Tagline directions: *"Keep every promise."* / *"The deals you win are the promises you keep."* Competitive framing: Gong tells you what was said; PromiseKeeper makes sure it happens — for both sides, inside Salesforce.

## 11. Competitive Landscape & Differentiation (VERIFIED — July 2026)

This section reflects live market research conducted in July 2026, not assumptions. Findings per category:

### 11.1 What actually exists today

**A. Conversation intelligence (closest overlap — take seriously).** Gong's Action Items feature has existed for years: it uses ML to detect action items in calls, lists them, emails reps, and logs them as Salesforce Task descriptions. Gong Engage now generates AI task recommendations based on commitment tracking, objection patterns, and buying signals,



and Gong's AI agents (AI Tasker, AI Deal Monitor) flag stalled deals and missed commitments. **Verified gaps:** action items are seller-oriented recommendations living primarily in Gong's ecosystem and pushed to CRM as flat one-sided Tasks; there is no bilateral buyer-obligation ledger, no structured promise lifecycle, no evidence-verified fulfillment, no external buyer nudging, and no reliability analytics computed from kept/broken history. Gong is also non-native (data leaves the Salesforce trust boundary) and priced for enterprise (~\$5K platform fee plus roughly \$1,360–\$1,600 per user annually) — leaving the mid-market wide open.

**B. Salesforce native (platform-risk zone).** Einstein Conversation Insights records and analyzes calls and automatically turns detected follow-ups into CRM Tasks tied to the opportunity, with next-step notifications for reps. And as of 2026, Slackbot's meeting intelligence can capture action items from meetings and update CRM records automatically. **Verified gaps:** everything Salesforce ships here is (a) call/meeting-centric, weak on email threads; (b) one-sided — it creates internal Tasks, never buyer obligations; (c) fire-and-forget — a Task is created, but nothing verifies fulfillment, detects breach, renegotiates, or computes reliability. Salesforce's own free "Action Plans" app is manual task templates. **Implication:** Salesforce is commoditizing *extraction*. Do not build the moat on extraction.

**C. Mutual action plans / digital sales rooms.** Dock, Recapped, Accord, SalesHood, and Topo.io (which is listed on AppExchange as a native digital sales room with mutual action plans) let sellers and buyers share plans with tasks assigned to either side. **Verified gaps:** every one of these requires the seller to *manually author* the plan; none extracts obligations automatically from real communication; they are shared workspaces (heavyweight, buyer must engage with a portal), not ambient ledgers; and none verifies fulfillment or tracks broken-promise history. They prove demand for bilateral accountability but solve it with manual labor.

**D. AI notetakers/meeting assistants (Fireflies, SUMO, Momentum, Oliv, etc.).** Commoditized: they transcribe, summarize, extract action items, and push tasks/notes into Salesforce. One-directional pipes into Tasks; no obligation object, no lifecycle, no bilateral view, no follow-through loop. These are potential *sources* for PromiseKeeper, not competitors — an integration-partner opportunity.

## 11.2 The verified white space

Mapping the market on two axes — **automatic vs. manual capture** and **one-sided tasks vs. bilateral verified obligations** — every existing player sits in three quadrants. The fourth is empty:

|                       | Manual capture                             | Automatic capture                            |
|-----------------------|--------------------------------------------|----------------------------------------------|
| One-sided tasks       | Salesforce Tasks, Action Plans             | Gong Action Items, ECI, notetakers, Slackbot |
| Bilateral obligations | Dock, Recapped, Accord, Topo (manual MAPs) | ← PromiseKeeper (empty quadrant)             |

And no player in *any* quadrant does the three things that constitute the real moat: (1) **evidence-verified fulfillment** (proving the promise was actually kept, not just marking a task done), (2) **breach/renegotiation lifecycle intelligence**, and (3) **reliability analytics** derived from the kept/broken record.

### 11.3 Differentiators — ranked by defensibility

1. **The Fulfillment Verification Engine (primary moat).** Closing the loop with evidence — matching the sent attachment, the booked meeting, the introduced contact — is the hardest capability to copy and the one nobody has. A Gong task says "send SOC 2"; PromiseKeeper *knows* it was sent, when, links the proof, and knows when it wasn't.
2. **The Promise object as native CRM data.** A structured, reportable, flow-automatable, API-accessible object with lifecycle and audit trail — inside the customer's org, inside the Einstein Trust Layer. Compounds into an irreplaceable dataset (switching cost) and is usable by the customer's own agents and flows (platform gravity).
3. **Bilateral accountability without a portal.** Buyer obligations tracked and courteously nudged through normal email — no deal-room login required. This captures the MAP category's value with zero buyer friction and zero rep authoring effort.
4. **The Reliability Graph (new — added in v2.0, see §6.8).** Kept/broken/renegeotiated history rolls up into reliability signals per contact, account, rep, team, and promise category. Novel data that literally cannot exist without the ledger — the analytics moat on top of the data moat.
5. **Commitment Escrow for conditional promises (new — see §6.9).** Structured tracking of "if X then Y" concessions — a chronic leak (especially sales→delivery handoffs) that no task, CI, or MAP product models at all.
6. **Post-sale continuity.** The ledger survives closed-won and flows into CS/delivery — CI tools stop at the sale; MAPs mostly stop at onboarding.
7. **Audit-grade obligation record.** Immutable Promise Events give regulated buyers (gov, healthcare, finserv) a defensible commitment trail — a compliance angle no adjacent product has.

### 11.4 Anti-commoditization strategy

Because extraction is commoditizing, PromiseKeeper must (a) treat ECI/notetaker/Slack outputs as *ingestible sources* (be the ledger of record above them, not a rival extractor); (b) lead every demo and marketing message with verification, bilateral ledger, and reliability

analytics — never "AI finds your action items"; (c) file the category flag early: "**Obligation Ledger / Commitment Intelligence**" as a distinct AppExchange positioning vs. "conversation intelligence" and "digital sales rooms."

### 11.5 Residual risks & final verification gate

- **Platform risk (highest):** Salesforce extending ECI/Slackbot from task creation to obligation tracking. Mitigations: speed, data moat, the verification engine (technically deep), post-sale breadth, and a partnership-friendly native posture.
  - **Gong down-market or deeper CRM writes.** Mitigations: native trust story, mid-market pricing, buyer-side + evidence capabilities Gong's seller-workflow DNA doesn't cover.
  - **Confidence in uniqueness (updated):** ~85% for the full concept as specified in v2.0 (up from 72% pre-verification — the differentiated surface is now larger and verified against real products). Residual uncertainty: unlisted/early-stage AgentExchange entrants.
  - **Mandatory gate before build kickoff:** a listing-by-listing AppExchange + AgentExchange sweep (search terms: commitment, promise, obligation, action item, mutual plan, follow-up agent, accountability) with screenshots archived — both as diligence and as prior-art documentation for positioning. Owner: PM. This PRD's concepts, object model, terminology, and positioning are original work product authored for this document; the sweep additionally protects against inadvertent similarity in naming or listing copy.
- 

## 12. Success Metrics

### 12.1 Product quality (model)

- Extraction precision  $\geq 90\%$ , recall  $\geq 85\%$  on evaluation set (explicit commitments, English, email) at GA; disputed-extraction rate  $< 5\%$  of confirmed promises.
- Due-date resolution accuracy  $\geq 90\%$  for explicit relative dates.
- Fulfillment auto-detection precision  $\geq 90\%$  (false "fulfilled" is worse than a missed detection).

### 12.2 Adoption & engagement

- Time-to-first-ledger  $< 10$  minutes from install;  $\geq 60\%$  weekly active among licensed reps by day 30;  $\geq 70\%$  of Proposed promises triaged within 48h; brief open rate  $\geq 50\%$ .

### 12.3 Customer outcomes (pilot targets)

- +30% on-time fulfillment of tracked commitments vs. pre-deployment baseline.
- Measurable cycle-time reduction on deals with active ledgers (target:  $-10\%$ ); qualitative NPS/testimonials on "walked into meetings prepared."

- Renewal cohort correlation study by quarter 2 of GA.

## 12.4 Business

- Free→paid conversion  $\geq 15\%$ ; net revenue retention  $\geq 120\%$  via seat + edition expansion; AppExchange rating  $\geq 4.7$ .
- 

## 13. MVP Scope & Roadmap

### 13.1 MVP (Months 0-4)

**In:** Email-only extraction (English); Promise object + lifecycle; record-page Ledger LWC; My Promises view; review queue; internal nudges + daily digest; manual + basic auto fulfillment (attachment/reply in thread); conversational agent (query + status commands); admin setup wizard; 3-user free edition mechanics; security-review submission prep. **Out (explicitly):** calls/meetings, external nudges, escalations, analytics dashboards, Slack, multilingual, shareable ledger. **MVP exit criteria:** quality bars in 12.1 on email corpus; 3 design-partner orgs live; 60-second demo flow validated.

### 13.2 Phase 2 (Months 5-8) — Premium launch

Call/meeting transcript sources; external nudge drafting with approval flow + tone guardrails; escalation policy builder; pre/post-meeting briefs; Promise Health Score + manager command center + report types; AppExchange GA.

### 13.3 Phase 3 (Months 9-14)

Slack app (nudges, briefs, ledger commands — fits AgentExchange Slack unification); **Commitment Escrow / conditional promises (§6.9); Reliability Graph v1 (§6.8);** source-agnostic ingestion connectors (ECI, Gong, Fireflies — §6.10); renegotiation intelligence; CRM Analytics dashboard pack; outcome-correlation reporting; multilingual v1; shareable external ledger (QBR mode).

### 13.4 Phase 4 (Months 15+) — Expansion bets

Post-sale SLA/contractual obligation tracking (CLM-adjacent SKU); partner-ecosystem three-way ledgers; promise-backed forecasting (§6.8/F-REL-3) feeding forecast models; benchmarks network (opt-in); API/Data Cloud packages; industry compliance packs (gov, healthcare) leveraging the audit-grade ledger.

---

## 14. Risks & Mitigations

| # | Risk                     | Likelihood | Impact | Mitigation                       |
|---|--------------------------|------------|--------|----------------------------------|
| 1 | Extraction quality below | Med        | High   | Verbatim-anchor rule, confidence |

| #  | Risk                                                             | Likelihood | Impact | Mitigation                                                                                                                    |
|----|------------------------------------------------------------------|------------|--------|-------------------------------------------------------------------------------------------------------------------------------|
|    | trust threshold → users disable                                  |            |        | routing, review queue, dispute feedback loop, published quality bar before GA                                                 |
| 2  | Notification fatigue → product ignored                           | Med        | High   | Digest-first design, batching, smart priority, aggressive default quiet settings                                              |
| 3  | External nudges damage a relationship                            | Low-Med    | High   | Approval-by-default, tone guardrails, frequency caps, sensitive-contact blackout, red-team suite                              |
| 4  | Reps feel surveilled ("promise tracking = performance tracking") | Med        | Med    | Rep-first value framing (briefs, cover-your-promises), manager views aggregate by default, configurable coaching-view rollout |
| 5  | Buyers feel surveilled if ledger sharing mishandled              | Low        | High   | Internal-only default; sharing is explicit, filtered, opt-in (Phase 3)                                                        |
| 6  | Platform risk: Salesforce ships native equivalent                | Med        | High   | Speed, data moat, bilateral + fulfillment depth, post-sale expansion; consider partnership posture                            |
| 7  | Prompt injection via inbound email                               | Med        | High   | Untrusted-data prompt architecture, schema-validated outputs, no tool execution from inbound content                          |
| 8  | Security review delays AppExchange listing                       | Med        | Med    | Native architecture, early submission, dedicated review workstream in MVP phase                                               |
| 9  | Email volume / governor limits in large orgs                     | Med        | Med    | Batch design, throttles, incremental backfill, Data Cloud path for scale                                                      |
| 10 | Privacy/works-council objections (EU) to mailbox processing      | Med        | Med    | Per-user opt-in modes, exclusion rules, minimized excerpts, DPA + regional guidance docs                                      |
| 11 | Overlap perception with Gong/MAPs blunts differentiation         | Med        | Med    | Category creation ("obligation ledger"), pre-build AppExchange sweep (11), demo that shows what CI can't do                   |

15. Open Questions

- 1. Backfill: how far back do we process historical email on install (value vs. cost vs. privacy)? Proposal: 90 days, admin-configurable.
- 2. Should soft/hedged commitments be tracked by default or off by default? (Design-partner test.)
- 3. Ownership routing: when the promiser isn't the fulfiller (rep promises, SE delivers) — auto-suggest owner by category, or always the promiser? (Proposal: suggest, human confirms.)
- 4. Closed-lost hygiene: auto-waive open promises at closed-lost, or preserve for re-engagement plays?
- 5. Pricing metric: pure per-seat vs. per-seat + promise-volume tier for very high-volume orgs?
- 6. Does the free demo edition process live data or offer a synthetic sandbox mode first (faster security comfort)?
- 7. Partnership vs. compete posture toward CI vendors (ingest Gong transcripts as a source = distribution wedge or dependency risk)?

16. Appendix

A. Example extraction (illustrative)

Email, Tue 14:02 — "Great call today. I'll get the SOC 2 report and the pen-test summary over to you by Friday. Could you send the current data schema and loop in whoever owns SSO on your side before our Thursday sync next week?"

| Promise                   | Direction   | Owner             | Due                             | Confidence |
|---------------------------|-------------|-------------------|---------------------------------|------------|
| Send SOC 2 report         | We owe them | Rep (suggest: SE) | Fri (this week)                 | Explicit   |
| Send pen-test summary     | We owe them | Rep (suggest: SE) | Fri (this week)                 | Explicit   |
| Share current data schema | They owe us | Contact: J. Doe   | Before Thu sync (resolved date) | Explicit   |
| Introduce SSO owner       | They owe us | Contact: J. Doe   | Before Thu sync (resolved date) | Explicit   |

B. Default promise categories

Documents & collateral · Security/compliance artifacts · Pricing & commercial ·

Legal/contract steps · Meetings & introductions · Technical follow-ups · Data/asset sharing · Approvals & internal escalations · Delivery/onboarding items · Support/fix commitments.

## C. Glossary

**Promise:** discrete trackable commitment with owner, beneficiary, due date, and source.

**Ledger:** bilateral set of promises on a record. **Promise debt:** open unfulfilled promises.

**Breach:** internal state for overdue-unfulfilled. **Nudge:** reminder instance, internal or external-draft. **Brief:** pre/post-meeting or daily ledger summary.

---

*End of document — PromiseKeeper PRD v2.0 (Market-Verified Edition)*